

¶ the evidence index — every claim on file · 11 entries

The evidence index

The master ledger behind every number on this site: the claim, the strongest artifact it terminates at, when it was recorded, and the case-file receipt that argues it in full. If a claim is not in this ledger or a case file, the site does not make it — apart from the biographical register (degree, dean’s list, GPA), which the transcript on request backs.

● ● ● ○ ● ● ● ● — ○ ○ ●

at a glance — 11 entries · 8 public · 3 private-safe — 1 of them held, not yet earned

e-01

● [public]

3-layer local classifier

Applied's classifier is a 3-layer rules, embeddings, and SetFit path – all three run on the desktop app and in the browser Space; the hosted web app runs the rules layer alone.

source: docs/ML_STRATEGY.md @ 36a2f54 ↗ (https://github.com/yadava5/applied/blob/36a2f54/docs/ML_STRATEGY.md)

date: 2026-07-26

verification: ml strategy doc read against the backend source at the pinned commit; the hosted limit is the cloud short-circuit in classifier/hybrid.py, argued at case-file receipt 09.

argued in: jobtracker case file · receipt 02 —

boundary: no private email content is shown.

e-02

● [public]

macro-F1 0.9791

Applied's rules and deterministic hybrid v3 gates both passed on the 96-sample eval set with macro-F1 0.9791.

source: baseline_hybrid_v3.json @ 36a2f54 ↗ (https://github.com/yadava5/applied/blob/36a2f54/backend/data/evaluation/baseline_hybrid_v3.json)

date: 2026-03-03

verification: committed 2026-03-03 baseline (deterministic profile) plus the public backend-ci gate run of 2026-04-20. the artifact is byte-identical at the new pin – re-read 2026-07-26, same 96 samples, same 0.9791.

argued in: jobtracker case file · receipt 05 —

boundary: the committed baseline json records metrics and label counts, not message content.

e-03

● [public]

271 backend tests

The Applied backend suite passed 271 tests locally, with 10 skipped, under the test/null-keyring environment.

source: backend/tests @ 36a2f54 ↗ (https://github.com/yadava5/applied/tree/36a2f54/backend/tests)

date: 2026-07-26

verification: local run against the pinned public test tree, 2026-07-26. the 10 skips are the postgres rls module, which needs a live database url and gets one from no workflow.

argued in: jobtracker case file · receipt 04 —

boundary: the suite runs with a null keyring; no private email or account data is involved.

e-04

● [private-safe]

7-phase AutoML lifecycle

Agentic AutoML presents a 7-phase ML workflow with LangGraph and MCP orchestration.

source: expo poster capture – private-safe —

[self-hosted – checked into this site's repository]

date: 2026-05

verification: private-safe senior design poster and local automl repository audit.

argued in: automl case file · receipt 02 —

boundary: uses demo/source-truth data and excludes private repository source.

e-05

● [public]

71 iOS tests

Visual Assist has audited XCTest model and utility coverage.

source: VisualAssistTests @ 22ebdaa
↗ (<https://github.com/yadava5/VisualAssist/tree/22ebdaa/VisualAssistTests>)

date: 2026-05

verification: function count audited in the public test tree at the pinned commit.

argued in: visual assist case file
· receipt 01 →

boundary: no live camera, location, or user sensor data is shown.

e-06

● [public]

1,145 automated tests

Cadence suite measured 2026-07:
634 frontend + 511 backend =
1,145 passing (vitest).

source: taskflow-calendar @ 69a59e7
↗ (<https://github.com/yadava5/taskflow-calendar/tree/69a59e7>)

date: 2026-07

verification: local vitest run against the pinned public source.

argued in: taskflow case file ·
receipt 01 →

boundary: no private data.

e-07

● [public]

3.5× openmp+simd dot kernel

The openmp+simd dot kernel runs 3.5× faster than the -O3 baseline (dot 256) in committed benchmarks.

source: BENCHMARKS.md @ c6e5c0b ↗ (<https://github.com/yadava5/fast-mnist-nn/blob/c6e5c0b/BENCHMARKS.md>)

date: 2025-12-26

verification: committed 2025-12-26 benchmark run data in the public fast-mnist-nn repository.

argued in: fast-mnist case file ·
receipt 02 →

boundary: no private data.

e-08

● [public]

~97% test accuracy

Glyph reaches ~97% test accuracy on MNIST after ~30 epochs — documented in the repo's README training notes; no committed eval artifact reproduces it yet.



held until a committed eval run earns it.

status: **held — not yet earned**
source: README.md @ c6e5c0b ↗ (<https://github.com/yadava5/fast-mnist-nn/blob/c6e5c0b/README.md>)
[self-authored — the author's own documentation]
date: 2026-07
verification: readme training notes checked against the source at the pinned commit. the number is documentation, not a committed eval run — it stays held until one is checked in.
argued in: fast-mnist case file · receipt 01 (held) →
boundary: no private data.

e-09

● [private-safe]

10,453 deduped rows

Master Inventory private-safe proof records the deduped inventory row count and schema boundary.

source: master-inventory-ledger.json — checked in 2026-06 →
[self-hosted — checked into this site's repository]
date: 2026-06
verification: sanitized processed-output ledger, checked into this site's public repository.
argued in: master inventory case file · receipt 01 →
boundary: raw institutional exports, report names, owner names, and rows are excluded.

e-10

● [private-safe]

19/20 structured sweep

PolicyBot private-safe
proof records a structured
validation sweep.

source: policybot-validation-
ledger.json — checked in 2026-06 →
[self-hosted — checked into this
site's repository]
date: 2026-06
verification: sanitized validation
ledger, checked into this site's
public repository.
argued in: policybot case file ·
receipt 01 →
boundary: raw policy text and slack
messages are excluded.

e-11

● [public]

72 tests pass

jetpack-compress compiles clean
on JDK 25 and its full suite
passes — 72 tests (Tests run: 72,
Failures: 0).

source: README.md @ af2c4b1 ↗ (<https://github.com/yadava5/jetpack-compress/blob/af2c4b1/README.md>)
[self-authored — the author's own
documentation]
date: 2026-07
verification: readme status line
read against the public repo at the
pinned commit; `mvn test` runs the
72-test junit 5 suite.
argued in: no case file — the
repository is the record
boundary: no private data.

the evidence index

back to the work ←

first case file — agentic automl →

© 2026 ayush yadav – set
by hand

set in fraunces, newsreader &
fragment mono

two inks. one line. seven chapters.