



ayush
yadav

the
work

experience

contact

github

motion:
on

 resume

¶ case file 01 / 07 · agentic automl platform – filed 2025-09 · last
verified 2026-07

status: in progress – core shipped, run ledger pending

you opened this file · jul 25, 2026

Agentic AutoML Platform

A private, GitHub-backed agentic AutoML platform. Datasets and domain documents become auditable pipeline decisions – and a human approval gate holds every generated action before it alters the workflow.

role capstone engineer – my slice below
timeframe 2025-09 to present
stack typescript · react 19 · express 5 · postgresql ·
docker · langgraph · mcp
repo private – yadava5/ai-augmented-auto-ml-toolchain
live demo agentic-automl-platform.vercel.app ↗ (https://agentic-automl-platform.vercel.app)



```
private proof: github evidence – the
current github repository is
yadava5/ai-augmented-auto-ml-
toolchain and its readme identifies
the product as agentic automl
platform. the repository is private,
so this case file shows private-safe
evidence instead of a public source
link.
```

[problem] · § as found

Between a raw dataset and a useful model sits a chain of repetitive judgment: ingestion, feature decisions, training, evaluation, deployment packaging. Automate the chain carelessly and the judgment disappears with the labor.

constraints –

make pipeline decisions auditable instead of opaque.

support domain documents through retrieval and mcp-based orchestration.

require human approval before generated actions alter workflow state.

keep training workflows reproducible with containerized execution.

validate the product flow with browser-level checks.




fig. 1 – the halt, echoed – the run crosses the 7-phase lifecycle and stops at the human gate.

description: A drawn figure that runs – not a screenshot. A card-scale echo of the pinned home scene: the ink run token crosses the lifecycle's first six phases and halts at the clay human gate – 7.0 deploy never resolves, because deploying is a person's decision, not this card's. The phases and the gate are the platform's real 7-phase, LangGraph + MCP-orchestrated workflow (default model GPT-5.4); run 041 stays awaiting its human.

A React and TypeScript interface coordinates Express/PostgreSQL services, LangGraph + MCP orchestration, notebook-based training workflows, Docker execution, and Playwright evaluation.



the circuit –

react 19 ui → express 5 api workflow requests
express 5 api → langgraph + mcp approved actions
express 5 api → docker runtime training jobs
docker runtime → postgresql 16 run records
playwright evals → react 19 ui browser proof & closes the loop

↗ generated actions hold at the approval edge
until a human says go

fig. 2 – the gated loop, inked. clay marks the gates: where a check can stop the run.

run	model	metrics	status
038	log-reg		approved
039	random forest		approved
040	gbm		approved
041	xgboost		awaiting approval

fig. 3 – experiment registry, transcribed private-safe excerpt.
 metrics withheld – private repository; see the boundary rows below.

[decisions] · § as filed

d1 – use langgraph + mcp for orchestration · accepted

The platform needs phase-aware routing, tool calls, and auditable decisions tied to domain context.^{t1}

t1. tradeoff – more infrastructure than a simple model runner, but better for traceable workflows.

d2 – keep human approval gates · accepted

Generated preprocessing, training, and deployment actions should be reviewed before they alter the workflow.^{t2}

t2. tradeoff – approval gates slow down full automation, but they make the system safer and easier to debug.

d3 – containerize execution · accepted

Training runs need reproducible environments.^{t3}

t3. tradeoff – docker adds setup cost but reduces machine-specific drift.

[validation] · § the receipts

method slip – eval protocol
protocol: not yet documented – see corrections.

validation

walk the 8 claims →

claim	method · date	artifact · visibility
01 The platform lives in the private repo yadava5/ai-augmented-auto-ml-toolchain; its README titles it Agentic AutoML Platform.	live github metadata check 2026-07	no linkable artifact – described only [private-safe]
02 Upload, EDA, NL-to-SQL, preprocessing, training, experiments, and deployment are the seven lifecycle phases.	read from the senior design poster and the local repo 2026-05	see fig. 6 – the expo poster [private-safe]
03 HPO, multi-model search, notebook training, and automated workflow evaluation are built into the platform.	poster + presenter deck, checked against the local repo 2026-05	see fig. 5 – the presenter deck [private-safe]
04 Presenter slide 8 records the stack and validation posture: all-green tests, coverage, logs, packages, and migrations.	transcribed from the presenter artifact 2026-05	see fig. 5 – the presenter deck [private-safe]

05 My slice of the build: the Monaco/Jupyter runtime with live WebSocket sync, Docker sandbox constraints, the eval runner, and the Optuna study streaming UI.	as presented – the presenter artifact names this work 2026-05	see fig. 5 – the presenter deck [private-safe]
06 Training runs execute in a Dockerized runtime for reproducibility.	poster architecture panel + local repo audit 2026-05	see fig. 6 – the expo poster [private-safe]
outcomes		
claim	method · date	artifact · visibility
07 A dataset and a goal become a structured, auditable workflow – planned and argued for by agents that still cannot press go.	the product's design, described – not an outcome metric date not recorded	no linkable artifact – described only [private-safe]
08 Pipeline decisions run through LangGraph and MCP tool calls rather than free-form output.	poster workflow-state panel 2026-05	see fig. 6 – the expo poster [private-safe]
ci rows link the public run · repo pins are the exact commits verified 2026-07.		

what i'm not claiming –

no per-run metrics are published here. the registry excerpt shows run, model, and status only – a demo-data run ledger with a complete metric trail has not shipped yet.

the repository is private, so source is not inspectable from this page. the evidence is the expo poster, the presenter deck, and the private-safe registry capture – verifiable in interview.

[corrections] · § the register

note · 2026-07

Per-run metrics remain withheld while the repository is private; the eval protocol for the platform's own model runs is not yet publicly documented. Nothing here has been retracted – this register is waiting on a demo-data run ledger.

[appendix] · § plates & artifacts




fig. 4 – private-safe experiment registry screenshot.
description: local automl repository, demo data · 2026-06
· open to inspect →







fig. 5 – presenter stack proof.
description: senior design presenter deck, slide 8 · 2026-05
· open to inspect →



§2 – an agent you can approve



§3 – one langgraph, every step auditable



§4 – measured on the public leaderboard

fig. 6 – expo poster proof.
panels cropped at column width
from the checked-in capture –
titles quoted from the poster.
description: senior design expo
poster, spring 2026 · 2026-05
open the full poster →

next file — ¶ 02 / 07 · glyph — filed 2025-10 →

© 2026 ayush yadav — set set in fraunces, newsreader &
by hand fragment mono

two inks. one line. seven chapters.