



ayush
yadav

the
work

experience

contact

github

motion:
on

 resume

¶ the evidence index — every claim on file · 11 entries

The evidence index

The master ledger behind every number on this site: the claim, the strongest artifact it terminates at, when it was recorded, and the case-file receipt that argues it in full. If a claim is not in this ledger or a case file, the site does not make it.

e-01

3-layer local classifier

Applied uses a 3-layer rules, embeddings, and SetFit classifier path for job-search email classification.

source: docs/ML_STRATEGY.md @
3225eb4 ↗ (https://github.com/yadava5/jobtracker/blob/3225eb4/docs/ML_STRATEGY.md)

date: 2026-06

verification: ml strategy doc read
against the backend source at the
pinned commit.

argued in: jobtracker case file ·
receipt 02 —

visibility: [public]

boundary: no private email content
is shown.

e-02

macro-f1 0.9791

Applied's rules and deterministic hybrid v3 gates both passed on the 96-sample eval set with macro-F1 0.9791.

source: baseline_hybrid_v3.json @ 3225eb4 ↗ (https://github.com/yadava5/jobtracker/blob/3225eb4/backend/data/evaluation/baseline_hybrid_v3.json)

date: 2026-03-03

verification: committed 2026-03-03 baseline (deterministic profile) plus the public backend-ci gate run.

argued in: jobtracker case file · receipt 05 —

visibility: [public]

boundary: the committed baseline json records metrics and label counts, not message content.

e-03

182 backend tests

The Applied backend suite passed 182 tests locally under the test/null-keyring environment.

source: backend/tests @ 3225eb4 ↗ (<https://github.com/yadava5/jobtracker/tree/3225eb4/backend/tests>)

date: 2026-06

verification: local run against the pinned public test tree.

argued in: jobtracker case file · receipt 04 —

visibility: [public]

boundary: the suite runs with a null keyring; no private email or account data is involved.

e-04

7-phase automl lifecycle

Agentic AutoML presents a 7-phase ML workflow with LangGraph and MCP orchestration.

source: expo poster capture – private-safe →
[self-hosted – checked into this site's repository]
date: 2026-05
verification: private-safe senior design poster and local automl repository audit.
argued in: automl case file · receipt 02 →
visibility: [private-safe]
boundary: uses demo/source-truth data and excludes private repository source.

e-05

71 ios tests

Visual Assist has audited XCTest model and utility coverage.

source: VisualAssistTests @ 22ebdaa
↗ (<https://github.com/yadava5/VisualAssist/tree/22ebdaa/VisualAssistTests>)
date: 2026-05
verification: function count audited in the public test tree at the pinned commit.
argued in: visual assist case file · receipt 01 →
visibility: [public]
boundary: no live camera, location, or user sensor data is shown.

e-06

1,145 automated tests

Cadence suite measured 2026-07: 634 frontend + 511 backend = 1,145 passing (vitest).

source: taskflow-calendar @ 69a59e7
↗ (<https://github.com/yadava5/taskflow-calendar/tree/69a59e7>)
date: 2026-07
verification: local vitest run against the pinned public source.
argued in: taskflow case file · receipt 01 →
visibility: [public]
boundary: no private data.

e-07

3.5x openmp+simd dot kernel

The openmp+simd dot kernel runs 3.5x faster than the -O3 baseline (dot 256) in committed benchmarks.

source: `BENCHMARKS.md @ c6e5c0b` ↗ (<https://github.com/yadava5/fast-mnist-nn/blob/c6e5c0b/BENCHMARKS.md>)

date: 2025-12-26

verification: committed 2025-12-26 benchmark run data in the public fast-mnist-nn repository.

argued in: fast-mnist case file · receipt 02 →

visibility: [public]

boundary: no private data.

e-08

~97% test accuracy

Glyph reaches ~97% test accuracy on MNIST after ~30 epochs — documented in the repo's README training notes; no committed eval artifact reproduces it yet.



held until a committed eval run earns it.

status: **held — not yet earned**

source: `README.md @ c6e5c0b` ↗ (<https://github.com/yadava5/fast-mnist-nn/blob/c6e5c0b/README.md>)

[self-authored — the author's own documentation]

date: 2026-07

verification: readme training notes checked against the source at the pinned commit. the number is documentation, not a committed eval run — it stays held until one is checked in.

argued in: fast-mnist case file · receipt 01 (held) →

visibility: [public]

boundary: no private data.

e-09
10,453 deduped rows
Master Inventory private-safe
proof records the deduped
inventory row count and schema
boundary.

source: master-inventory-
ledger.json — checked in 2026-06 →
[self-hosted — checked into this
site's repository]
date: 2026-06
verification: sanitized processed-
output ledger, checked into this
site's public repository.
argued in: master inventory case
file · receipt 01 →
visibility: [private-safe]
boundary: raw institutional
exports, report names, owner names,
and rows are excluded.

e-10
19/20 structured sweep
PolicyBot private-safe proof
records a structured validation
sweep.

source: policybot-validation-
ledger.json — checked in 2026-06 →
[self-hosted — checked into this
site's repository]
date: 2026-06
verification: sanitized validation
ledger, checked into this site's
public repository.
argued in: policybot case file ·
receipt 01 →
visibility: [private-safe]
boundary: raw policy text and slack
messages are excluded.

e-11

72 tests pass

jetpack-compress compiles clean
on JDK 25 and its full suite passes
– 72 tests (Tests run: 72, Failures:
0).

source: README.md @ af2c4b1 ↗ (<https://github.com/yadava5/jetpack-compress/blob/af2c4b1/README.md>)

[self-authored – the author's own
documentation]

date: 2026-07

verification: readme status line
read against the public repo at the
pinned commit; `mvn test` runs the
72-test junit 5 suite.

argued in: no case file – the
repository is the record

visibility: [public]

boundary: no private data.

the evidence index

back to the work ←

first case file – agentic automl →

© 2026 ayush yadav – set
by hand

set in fraunces, newsreader &
fragment mono

two inks. one line. seven chapters.